

# Unit-4

## LINEAR DISCRIMINANTS FOR MACHINE LEARNING

### INTRODUCTION TO LINEAR DISCRIMINATES

#### What is LDA (Linear Discriminant Analysis)?

LDA is a **classification** and **dimensionality reduction** technique that tries to **find a new axis (or direction)** that separates different classes of data as clearly as possible.

Think of it as drawing a straight line (in 2D) or a plane (in higher dimensions) that best splits the categories apart.

---

#### Main Goal of LDA

- **Maximize the distance** between different classes (class separation)
- **Minimize the spread** of data points within the same class

This means LDA tries to **keep classes apart** and **keep each class tight** and compact.

---

#### How LDA Works — Step-by-Step

Let's say we have data from two types of flowers: **Class A** and **Class B**, and each flower is described by features like petal length and width.

##### 1. Compute the mean (average) of each class

- Find the center point (mean) of Class A and Class B.

##### 2. Compute the scatter (spread) within each class

- Measure how spread out the data points are within each class.

##### 3. Compute the scatter between the classes

- Measure how far apart the means of the classes are.

##### 4. Create a linear combination of features (a line)

- LDA finds a new axis (direction) where:

- The **classes are far apart (large between-class scatter)**
- The **points within the same class are close together (small within-class scatter)**

This new axis is where the classes are easiest to separate.

---

### Simple Example

Imagine plotting students based on two features:

- Hours studied
- Number of classes skipped

And we want to classify them as **Pass** or **Fail**.

LDA will find the best line (a weighted combo of "hours studied" and "classes skipped") that:

- Pushes the **Pass** and **Fail** students as far apart as possible.
- Keeps students within each group close together.

So instead of using both features separately, we use **one new feature** (along that line), making classification easier and faster.

---

### Difference from PCA

Feature	LDA	PCA
Type	Supervised (uses class labels)	Unsupervised (ignores labels)
Goal	Maximize class separability	Maximize overall variance
Output	Best axis for classification	Best axis for compression

---

### When to Use LDA

- You have **labeled data** (e.g., “spam” vs “not spam”)
- You want to **reduce dimensions** but keep class differences intact
- You want to **build a classifier** (e.g., to predict categories)

## LINEAR DISCRIMINATES FOR CLASSIFICATION

## 🧠 What is LDA really doing in classification?

LDA is like a smart separator. Imagine you have two kinds of fruits—**apples** and **oranges**—and you measure their **color** and **weight**.

Now you want to:

- Draw a line (or curve) to **separate apples from oranges**
- And when you see a new fruit, **decide which group it belongs to**

That's **classification**, and LDA is a mathematical way to **find that best separating line**.

---

## 🏠 LDA Classification: Intuition

Suppose:

- You have **two classes**: A and B (for example, spam and not spam)
- Each class has some **data points** (samples)
- Each data point has **features** (like word counts in spam emails)

LDA will:

1. Look at how **spread out** each class is (within-class scatter)
2. Look at how **far apart** the classes are (between-class scatter)
3. Find a **new axis** (a straight line) that keeps:
  - The **classes far apart**
  - The **data in each class tightly packed**

Once this line is found, it projects the data onto this line and finds a **threshold (cutoff point)** to decide which side is class A and which side is class B.

---

## 📐 Mathematical Insight (Simplified)

LDA builds a function like:

$$y = w^T x + w_0$$

Where:

- $x$  is the input (like feature values)

- $w$  is the direction vector (how the features are combined)
- $w_0$  is a bias/threshold term
- $y$  is the result — if it's closer to Class A's mean, classify as A; else, B.

So LDA turns your **multi-dimensional data** into a **1D number**, and based on that number, it makes a decision.

---

### 🔧 Why LDA Works Well

- It **considers class labels** when reducing dimensions.
  - It **uses probability** and assumes each class has a **normal distribution**.
  - It works well when your classes are **linearly separable** (can be separated by a straight line or plane).
- 

### 🔍 Visual Example (Imagine)

Let's say we plot students' performance on a 2D graph:

- X-axis = study time
- Y-axis = sleep time

Red dots = "Pass", Blue dots = "Fail"

The red and blue dots form two cloud-like shapes. LDA tries to:

- Find the best straight line (diagonal maybe)
  - When you project all points onto this line, Pass and Fail should become well-separated along it
  - Then you can draw a cutoff point on that line and classify new students
- 

### 🤔 How is this better than just drawing any line?

Other methods might try to guess a line just by trial and error. LDA **uses math to optimize**:

- Maximize **between-class distance**
- Minimize **within-class scatter**

This means it tries to make the classes **as far from each other** and **as compact as possible**.

---

## Summary – LDA for Classification

Element	Description
<b>Goal</b>	Separate classes clearly using a linear boundary
<b>Works Best When</b>	Classes are normally distributed and have similar variances
<b>How it Decides</b>	Projects data onto a line and assigns to the nearest class
<b>Output</b>	A class label for each new input (e.g., spam or not spam)

### PERCEPTION CLASSIFIER:

## 1. What is a Perceptron, REALLY?

A **Perceptron** is the **simplest form of an artificial neuron**.

It:

- Takes multiple **inputs**
- Multiplies each input by a **weight**
- Adds a **bias**
- Runs the result through an **activation function**
- Makes a **binary decision** (Class 0 or Class 1)

Think of it like a **yes/no switch** — it turns ON (1) if the input is strong enough, OFF (0) if not.

---

## 2. Perceptron Components in Detail

### Inputs (x)

A list of numbers representing features. For example, for a flower:

- $x_1$  = petal length
- $x_2$  = petal width

### Weights (w)

Each feature is multiplied by a weight to show how important it is. For instance:

- If petal width is more important, its weight might be higher.

### Bias (b)

A constant added to shift the decision boundary — it helps the model make better decisions.

### Sum and Activation

All weighted inputs are added up:

$$z = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

Then, we apply an **activation function** — in the original perceptron, it's just:

$$\text{output} = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{otherwise} \end{cases}$$

### 3. Geometric View

Imagine your data on a 2D graph — red and blue dots.

The perceptron tries to find a **straight line** that divides the two:

- One side = Class 0
- Other side = Class 1

That line is based on the weights and bias:

$$w_1x_1 + w_2x_2 + b = 0$$

In higher dimensions (3D, 4D, etc.), this becomes a **hyperplane**.

---

### 4. Learning Algorithm: Step-by-Step

Let's say we have some **training examples** — inputs with correct labels.

The perceptron learns like this:

1. Start with **random weights**
2. For each example:
  - Compute predicted output
  - Compare it to the **true label**
  - If wrong, **adjust the weights**:

$$w_i = w_i + \eta(y_{\text{true}} - y_{\text{pred}})x_i$$

$$b = b + \eta(y_{\text{true}} - y_{\text{pred}})$$

Where:

- $\eta$  is the **learning rate** (e.g., 0.01)
- $y_{\text{true}}$  : real answer
- $y_{\text{pred}}$  : model's guess

It keeps doing this until it makes no more mistakes or reaches a maximum number of iterations.

## 🧠 5. Why It Works (And When It Doesn't)

The Perceptron works great when the data is **linearly separable**, meaning:

You can draw a straight line (or hyperplane) between the two classes without any mistakes.

BUT:

- If the data is **not linearly separable** (like XOR logic), it will **never converge** — it will keep looping and adjusting forever.

That's where **Multi-Layer Perceptrons (MLP)** or **deep neural networks** come in.

## 🔗 6. Real-World Analogy

Imagine you're the guard at a VIP club. You decide who enters based on:

- What kind of clothes they wear
- Whether they have a VIP card

You assign **weights** to those features:

- VIP card = very important (high weight)
- Clothes = kind of important (lower weight)

Then you:

- Add it up
- Apply a rule (if the total score  $> 0$ , let them in)

That's how the perceptron decides **yes or no**.

---

## 7. Real Example with Data

**Age Hours Studied Pass (1) / Fail (0)**

18	2	0
22	8	1
20	4	0
23	10	1

The perceptron will try to learn weights for:

- Age
- Hours studied

And learn the rule to classify who is likely to **pass or fail**.

---

## ☒ 8. Summary Table

Concept	Value
Input	Feature vector
Output	Class 0 or 1
Model Type	Linear binary classifier
Decision Surface	Hyperplane
Activation	Step function
Learns By	Updating weights if it makes errors
Good For	Simple classification tasks
Limitation	Can't handle complex or curved boundaries



## PERCEPTRON LEARNING ALGORITHM

### ◇ What is a Perceptron?

The **Perceptron** is a type of **artificial neuron** and the foundation of neural networks. It's a **supervised learning algorithm** used for **binary classification**—i.e., classifying data into two groups such as *yes/no*, *spam/not spam*, *positive/negative*, etc.

It learns to classify data by finding a **linear decision boundary (a straight line, plane, or hyperplane)** that separates the input data into two categories.

---

### ◇ Structure of a Perceptron

A perceptron consists of:

- **Inputs:** Features or attributes of the data (e.g., weather conditions, email content).
- **Weights:** Each input has a corresponding weight that signifies its importance.
- **Bias:** A constant value added to the weighted sum, helping shift the decision boundary.
- **Activation Function:** Determines the output — usually a **sign function** that outputs +1 or -1.

**Mathematical Expression:**

$$y = \text{sign}(w \cdot x + b)$$

Where:

- $x$  = input vector (features)
- $w$  = weight vector
- $b$  = bias
- $y$  = output (either +1 or -1)
- $\text{sign}$  = activation function that outputs 1 if the result is positive, else -1

---

### ◇ Perceptron Learning Rule

The perceptron learns by **adjusting the weights** based on its prediction errors. When a prediction is wrong, it updates the weights so it's less likely to make the same mistake again.

**Learning Rule (Update Equations):** If the prediction is incorrect:

$$w = w + \eta \cdot y \cdot x$$

$$b = b + \eta \cdot y$$

Where:

- $\eta$  = learning rate (a small constant, e.g., 0.1)
  - $y$  = true label (+1 or -1)
  - $x$  = input vector
- 

#### ◇ Steps of the Perceptron Learning Algorithm

1. **Initialize** weights and bias (usually to zero or small random values).
  2. For each input sample:
    - Calculate output using the current weights and bias.
    - If the output is correct, do nothing.
    - If incorrect, **update the weights and bias** using the learning rule.
  3. Repeat the process for all training samples.
  4. Continue iterating until:
    - All samples are correctly classified, or
    - A maximum number of iterations is reached.
- 

#### ◇ Conditions for Convergence

- The Perceptron algorithm **guarantees convergence** if the data is **linearly separable**, meaning a straight line (or hyperplane) can perfectly divide the two classes.
  - If the data is **not linearly separable**, the algorithm will **never converge**.
- 

#### ◇ Real-Life Example: Email Spam Detection

Imagine you want to classify emails as **spam (1)** or **not spam (-1)** using the presence of certain features:

- $x_1$ : Contains the word “Free”
- $x_2$  : Contains a link
- $x_3$  : Subject in ALL CAPS

The perceptron takes these inputs, multiplies them by weights, adds them up, and decides if the email is spam or not. Over time, it learns the correct pattern by adjusting the weights based on mistakes — just like a human adjusting their thinking after being corrected.

---

#### ◇ Advantages

- Simple and fast
- Easy to implement
- Good for linearly separable problems
- Forms the foundation of more complex neural networks

---

#### ◇ Limitations

- Works **only for binary classification**
- Can’t solve **non-linearly separable problems** (e.g., XOR problem)
- Doesn’t work well on **noisy data**
- Not suitable for **multi-class** classification unless extended

---

#### 📌 Summary

Feature	Description
Type	Supervised, binary classification
Model	Linear
Learning	By updating weights using errors
Converges when	Data is linearly separable
Output	+1 or -1

## Support vector machines

Svm stands for “support vector machines”

### 1. Core Idea Behind SVM

SVM tries to **find the best boundary** (hyperplane) that divides data into two classes.

The **best boundary** means:

- Not just any line separating the classes
- But the one with the **maximum margin**, i.e., the largest distance between the boundary and the nearest data points.

**Why maximum margin?**

- Larger margins mean better generalization on unseen data.
  - If the margin is small, small errors or noise could make wrong predictions.
- 

### 2. What is a Hyperplane?

A **hyperplane** is just a fancy term for a decision boundary.

- In 2D → hyperplane = **line**
- In 3D → hyperplane = **plane**
- In more dimensions → still called a **hyperplane**

A hyperplane separates the space into two halves. SVM tries to find the best one that cleanly separates the classes.

---

### 3. Support Vectors

Support Vectors are the **critical data points** closest to the hyperplane.

- These points “**support**” or define the position of the hyperplane.
- If you removed them, the boundary would change!
- They are the most important elements in training SVM.

**Visualize it like this:**

Imagine a stretched tight rope held by two poles. The poles closest to the rope are like the **support vectors**—they control where the rope is tied.

---

#### 4. SVM in Linearly Separable Case

When two classes can be separated perfectly by a straight line (without any overlap or noise):

- SVM finds the hyperplane with **maximum margin**.
- It ensures that:

$$y_i(w \cdot x_i + b) \geq 1$$

for every point:  $(x_i, y_i)$

Where:

- $w$  = weight vector (direction of the hyperplane)
- $b$  = bias (offset of the hyperplane)
- $y_i$  = class label (+1 or -1)

---

#### 5. Soft Margin SVM: Handling Imperfect Data

Real-world data is often **not clean**:

- Some points overlap.
- Some points are mislabeled.

**Soft Margin SVM** allows a few points to be on the wrong side of the margin or even the hyperplane.

It introduces a **penalty parameter C**:

- **High C**: Try hard to classify every point correctly (less tolerance for error).
- **Low C**: Allow more mistakes for a bigger margin (more flexible).

---

#### 6. When Data is Not Linearly Separable

Some datasets cannot be separated by a straight line, no matter what.

**Example:** XOR pattern:

**x1 x2 Output**

0 0 0

0 1 1

1 0 1

1 1 0

No straight line can separate this perfectly.

**Solution: Use the Kernel Trick!**

- **Kernel functions** allow SVM to project data into a **higher-dimensional space**.
- In that higher space, the data can **become linearly separable**.

Think of it like:

- In 2D, circles can't be separated by a line.
- But if you lift the circles into 3D, you can separate them with a plane!

---

## 7. Kernel Functions

**Kernel** = mathematical trick to compute in higher dimensions without explicitly transforming the data.

Popular kernels:

- **Linear Kernel:** Good for simple problems.
- **Polynomial Kernel:** Good for curved decision boundaries.
- **Radial Basis Function (RBF) or Gaussian Kernel:** Good for very complex boundaries.

RBF kernel formula:

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$$

- $\gamma$  = controls how far the influence of a single training example reaches.

---

## 8. Advantages of SVM

Advantage	Why it Matters
High-dimensional effectiveness	Works well even with a lot of features
Robust margin	Good generalization to new, unseen data
Versatile	Works with different kernels for non-linear problems
Few support vectors	Model complexity depends on support vectors, not dataset size

---

## 9. Disadvantages of SVM

Limitation	Problem
Computationally expensive	Slow for very large datasets
Hard to choose kernel and parameters	Needs experience or cross-validation
Not good for overlapping classes	When classes mix a lot, SVM struggles

---

## 10. Real-Life Applications of SVM

- **Face Detection** (classify part of an image as face or not)
- **Spam Detection** (classify emails as spam or not)
- **Image Classification** (cat vs dog, etc.)
- **Bioinformatics** (classifying genes, proteins)
- **Handwriting Recognition** (digit classification 0–9)

## LINEAR NON-SEPARABLE CASE NON LINEAR SVM

### ❧ 1. Why Some Data is Linearly Non-Separable

Not all real-world data is clean and neatly separable by a straight line or flat plane.

**Linearly non-separable** means:

- The two classes **overlap** in the feature space.
- No single straight line (or hyperplane) can cleanly separate them.

### ✂ Examples in real life:

- Handwritten digit "8" and "3" — sometimes parts overlap.
- Disease vs. healthy data — some symptoms are shared.
- Spam vs. not-spam — certain words appear in both.

---

## 🧠 2. What SVM Does to Handle This

SVM solves this in **two major ways**:

**A. Soft Margin – Tolerate some mistakes**

**B. Kernel Trick – Map data to higher dimensions**

Let's look at each.

---

### ♦ A. Soft Margin SVM – When Perfect Separation is Not Possible

SVM introduces **slack variables**  $\xi_i$  to allow for **violations** of the margin:

- Some points can be inside the margin.
- Some may even be on the wrong side of the hyperplane.

This leads to a new **optimization objective**:

$$\min \left( \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i \right)$$

Subject to:

$$y_i(w \cdot x_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$$

Where:

- $\xi_i$ : slack variables (0 if correctly classified with margin)
- $C$ : regularization parameter (penalty for misclassification)

**Intuition:**

- Large  $C$ : Tries to fit all points correctly (less tolerant of errors)
- Small  $c$ : Allows more margin violations (more tolerant)



This flexibility improves **generalization** on noisy or overlapping data.

### ♦ B. Kernel Trick – Make the Problem Linearly Separable in Higher Space

If the data can't be separated by a **line** in the original space, maybe it **can be** in a higher-dimensional space.

#### Idea:

Map original data  $x$  into a new feature space  $\phi(x)$ , so that in the new space, it **is** linearly separable.

But calculating  $\phi(x)$  directly can be computationally expensive or even infinite-dimensional!

So instead, we use the **kernel trick**:

$$K(x_i, x_j) = \phi(x_i) \cdot \phi(x_j)$$

This calculates the dot product **in the new space** without explicitly computing the mapping.



## 3. Common Kernel Functions – Transforming the Feature Space

Kernel Name	Formula	Use Case
Linear	$K(x, x') = x \cdot x'$	For linearly separable or almost-separable data
Polynomial	$K(x, x') = (x \cdot x' + 1)^d$	For curves and polynomial boundaries
RBF / Gaussian	$K(x, x') = \exp(-\gamma \ x - x'\ ^2)$	For complex, non-linear boundaries
Sigmoid	$\tanh(\alpha x \cdot x' + c)$	Inspired by neural networks



**Most popular: RBF (Radial Basis Function)** — powerful and works in many real-world scenarios.



## 4. Geometric Intuition of the Kernel Trick

Imagine:

- You have two concentric circles (like a donut).
- Class A = inside circle

- Class B = outside ring

No straight line in 2D can separate them.

But if you **lift the inner circle** upward into 3D (e.g., based on distance from origin), you can separate the classes using a **flat plane** in that 3D space.

SVM uses **kernel functions** to achieve this "lifting" mathematically.

## ✂ 5. Final SVM Optimization (With Kernel)

The SVM now solves:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j K(x_i, x_j)$$

Subject to:

- $0 \leq \alpha_i \leq C$
- $\sum_i \alpha_i y_i = 0$

Here:

- $\alpha_i$  : Lagrange multipliers (found via optimization)
- $K(x_i, x_j)$  : Kernel function (mapping via dot product)

Once you solve for  $\alpha_i$ , the decision function becomes:

$$f(x) = \sum_i \alpha_i y_i K(x_i, x) + b$$

## 🌀 6. Real-World Analogy

Imagine trying to separate black and white marbles:

- They are scattered on a table (2D), with no clear straight boundary.
- You add a **third dimension** by assigning height based on color features.
- Suddenly, they **separate vertically**, and you can use a flat sheet (hyperplane) to split them.

That's the essence of **kernel SVM**.

## ☑ 7. Summary – Key Takeaways

Concept	Description
Linearly Non-Separable	Classes cannot be separated with a straight line
Soft Margin	Allows some classification errors with a trade-off (C parameter)
Kernel Trick	Projects data into higher space for separation
Common Kernels	RBF (most used), Polynomial, Linear
Result	Non-linear boundary in original space becomes linear in mapped space
Real Impact	Works on complex problems like face recognition, spam detection, bioinformatics

### KERNAL – TRICK

The **kernel trick** is a fundamental concept in **Support Vector Machines (SVM)** and other machine learning algorithms. It allows us to solve complex problems by transforming the input data into a higher-dimensional space **without actually computing the transformation explicitly**. This enables SVMs to classify data that is not linearly separable in its original space.

---

### 🔍 Why Do We Need the Kernel Trick?

SVMs are **linear classifiers** by default. They find a hyperplane that best separates the classes. But what if the data isn't linearly separable?

Example:

Imagine trying to separate two classes shaped like concentric circles. A straight line won't work. But if we could somehow **transform** the data into a higher dimension, where the classes become linearly separable, we could apply a linear SVM there.

---

## 💡 What Is the Kernel Trick?

The kernel trick is a **mathematical shortcut**. It avoids the **explicit mapping** of data into higher dimensions. Instead, it uses a **kernel function** that computes the **dot product** of the data in the transformed space **directly from the original input space**.

Mathematically:

- Suppose  $\phi(x)$  maps input  $x$  into a higher-dimensional space.
- Instead of computing  $\phi(x)$ , we use a kernel function:

$$K(x, y) = \phi(x)^T \phi(y)$$

So, we compute the inner product in the transformed space **without ever computing  $\phi(x)$** . This is the **kernel trick**.

---

## 🔑 Common Kernel Functions

### 1. Linear Kernel:

$$K(x, y) = x^T y$$

- (No transformation; same as original space)

### 2. Polynomial Kernel:

$$K(x, y) = (x^T y + c)^d$$

- (Transforms into polynomial space of degree  $d$ )

### 3. Radial Basis Function (RBF) or Gaussian Kernel:

$$K(x, y) = \exp \left( -\frac{\|x - y\|^2}{2\sigma^2} \right)$$

- (Maps to infinite-dimensional space; great for non-linear problems)

### 4. Sigmoid Kernel:

$$K(x, y) = \tanh(\alpha x^T y + c)$$

- (Related to neural networks)

## ☑ Benefits of the Kernel Trick

- Handles **non-linearly separable data** effectively.
- Avoids the **computational burden** of working in high-dimensional spaces.
- Enables SVMs to model **complex decision boundaries**.

---

## 🔗 Real-World Analogy

Think of trying to separate red and blue marbles on a sheet of paper where they form circular patterns. You can't draw a line to separate them on the paper. But if you **lift** the paper (move to 3D), you might be able to separate them using a plane. The kernel trick is like "lifting" the data to a higher space **without actually lifting it**—you just know how far apart the marbles are in that lifted space using a special formula (kernel function).

## LOGISTIC REGRESSION:

Logistic regression is a statistical method used for binary classification, where the outcome variable has two possible categories (e.g., ☺ / ☹ / Ⓢ / Ⓣ). Instead of predicting a continuous outcome like in linear regression, logistic regression models the probability of the dependent variable belonging to a particular category. 🧠

Here's a breakdown of key aspects:

### How it Works:

1. **The Logistic Function (Sigmoid):** Logistic regression uses a sigmoid function (S-shaped curve ☒) to map the linear combination of input features to a probability value between 0 and 1. The formula for the sigmoid function is:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

where  $z$  is the linear combination of the independent variables:

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k$$

1. Here,  $\beta_0$  is the intercept,  $\beta_1, \beta_2, \dots, \beta_k$  are the coefficients for the independent variables  $x_1, x_2, \dots, x_k$ . **+ - X**
2. **Modeling Probability:** The output of the logistic function,  $\sigma(z)$ , is interpreted as the probability of the outcome being in the positive class (usually coded as 1). For example, if  $\sigma(z)=0.8$ , it means there's an 80% chance 🌀 that the outcome is 1.
3. **Decision Boundary:** To make a classification, a threshold (often 0.5) is set. If the predicted probability is above the threshold, the instance is classified into one class ☒; otherwise, it's classified into the other **X**.
4. **Parameter Estimation:** The coefficients ( $\beta_0, \beta_1, \dots, \beta_k$ ) are estimated using methods like Maximum Likelihood Estimation (MLE) to find the values that best fit the observed data. ⚙️

### Assumptions of Logistic Regression:

While logistic regression doesn't have as many strict assumptions as linear regression, some key assumptions include:

- **Binary Outcome:** The dependent variable must be binary (two categories). 🌀
- **Independence of Observations:** Each data point should be independent of the others. 🧑 🧑 🧑
- **Linearity of Log-Odds:** The relationship between the independent variables and the log-odds of the outcome is assumed to be linear. The log-odds (logit) is given by  $\ln(1-p)$ , where  $p$  is the probability of the event. 📏
- **No Multicollinearity:** The independent variables should not be highly correlated with each other. 🧑 🧑
- **Sufficiently Large Sample Size:** A larger sample size generally leads to more reliable estimates. 🧑
- **No Outliers:** Outliers can disproportionately influence the model. ⚠️

### Applications of Logistic Regression:

Logistic regression is widely used in various fields for binary classification tasks, including:

- **Healthcare:** Predicting the likelihood of a disease (e.g., diabetes, ❤️ heart disease) based on patient characteristics.
- **Finance:** 💰 Credit scoring (predicting loan default 🚫), fraud detection 🕵️.

- **Marketing:** 📈 Customer churn prediction 📄 ➡️, predicting the likelihood of a customer purchasing a product 🛒.
- **Natural Language Processing (NLP):** 🗣️ Sentiment analysis (😊/😞), spam detection 📧 ➡️ 🚫.
- **Ecology:** 🌳 Predicting the presence or absence of a species in a given habitat 🐾.
- **Social Sciences:** 👥 Predicting voting behavior 🗳️, participation in programs.

### Types of Logistic Regression:

- **Binary Logistic Regression:** The dependent variable has two categories. This is the most common type. 👉
- **Multinomial Logistic Regression:** The dependent variable has three or more unordered categories (e.g., predicting the type of movie a person will watch 🎬 🎭).
- **Ordinal Logistic Regression:** The dependent variable has three or more ordered categories (e.g., customer satisfaction levels: low 😞, medium 😐, high 😊).

### LINEAR REGRESSION MULTI-LAYER PERCEPTRONS(MLP) :

#### Linear Regression 📐

**Concept:** Linear regression is a relatively simple yet powerful algorithm used for predicting a continuous numerical output. It assumes a linear relationship between the independent variables (features) and the dependent variable (target). The goal is to find the best-fitting straight line (in the case of one independent variable) or hyperplane (in the case of multiple independent variables) that minimizes the difference between the predicted values and the actual values.

#### Mathematical Representation:

For a single independent variable (x) and a dependent variable (y), the linear regression model is represented as:

$$y = \beta_0 + \beta_1 x + \epsilon$$

where:

- y is the predicted value of the dependent variable.
- $\beta_0$  is the intercept (the value of y when x is 0).
- $\beta_1$  is the coefficient (slope) that represents the change in y for a one-unit change in x.
- $\epsilon$  is the error term, representing the difference between the actual and predicted values.

For multiple independent variables ( $x_1, x_2, \dots, x_n$ ), the equation becomes:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

### Key Characteristics:

- **Linearity:** Assumes a linear relationship between variables.
- **Simplicity:** Easy to understand and interpret.
- **Closed-form solution:** In many cases, the optimal parameters can be calculated directly.
- **Limited complexity:** Cannot model complex, non-linear relationships in the data effectively.
- **Applications:** Predicting house prices, sales forecasting, stock price prediction (in some cases), and other tasks where a linear relationship is expected.

### Multi-Layer Perceptron (MLP) 🧠

**Concept:** A Multi-Layer Perceptron (MLP) is a type of feedforward artificial neural network. It consists of multiple layers of interconnected nodes (neurons): an input layer, one or more hidden layers, and an output layer. Each neuron in one layer connects to every neuron in the subsequent layer. The connections have associated weights, and neurons typically have a non-linear activation function. This architecture allows MLPs to learn complex, non-linear patterns in data.

### Architecture:

- **Input Layer:** Receives the input features. The number of neurons corresponds to the number of features.
- **Hidden Layers:** One or more layers between the input and output layers. These layers perform non-linear transformations of the input data. The number of hidden layers and the number of neurons in each layer are hyperparameters that can be tuned.
- **Output Layer:** Produces the final prediction. The number of neurons and the activation function in this layer depend on the task (e.g., one neuron for regression, multiple neurons with SoftMax for multi-class classification).
- **Weights:** Determine the strength of the connections between neurons. These are learned during the training process.

### Key Characteristics:

- **Non-linearity:** Can model complex, non-linear relationships in the data.
- **Flexibility:** The architecture (number of layers, neurons) can be adjusted to suit the complexity of the problem.



- **Feature learning:** Can automatically learn relevant features from the input data through its hidden layers.
- **Computational cost:** Training can be computationally expensive, especially for deep networks with large datasets.
- **Interpretability:** Can be more difficult to interpret the learned relationships compared to linear regression.
- **Applications:** Image recognition, natural language processing, speech recognition, complex classification and regression tasks.

### Key Differences Summarized

Feature	Linear Regression	Multi-Layer Perceptron (MLP)
Relationship	Assumes a linear relationship	Can model non-linear relationships
Complexity	Simple	Complex
Layers	Single layer (input to output)	Multiple layers (input, hidden, output)
Activation	No non-linear activation function	Non-linear activation functions in neurons
Feature Learning	Requires manual feature engineering	Can automatically learn features
Interpretability	Generally easy to interpret coefficients	Can be difficult to interpret
Computational Cost	Relatively low	Can be high, especially for deep networks
Applications	Linear modeling tasks	Complex, non-linear tasks

# BACK PROPAGATION FOR TRAINING AN (MLP) :

**Backpropagation** is the core algorithm used to train Multi-Layer Perceptron (MLP) neural networks. It's an efficient way to calculate the gradients of the loss function with respect to the network's weights. These gradients are then used by optimization algorithms like Gradient Descent (or its variants) to update the weights and minimize the error of the network's predictions.

Here's a breakdown of how backpropagation works in training an MLP:

## 1. Forward Pass:

- An input data point is fed into the network.
- The input propagates through the network, layer by layer.
- At each neuron in the hidden and output layers:
  - The weighted sum of its inputs (from the previous layer) is calculated:

$$z_j = \sum_i w_{ij} a_i + b_j$$

- where:
  - $z_j$  is the weighted sum of inputs to neuron  $j$ .
  - $w_{ij}$  is the weight of the connection from neuron  $i$  in the previous layer to neuron  $j$  in the current layer.
  - $a_i$  is the activation of neuron  $i$  in the previous layer.
  - $b_j$  is the bias of neuron  $j$ .
- An activation function ( $\sigma$ ) is applied to the weighted sum to produce the output of the neuron:

$$a_j = \sigma(z_j)$$

Common activation functions include sigmoid, ReLU, and tanh.

- This process continues until the output layer produces a prediction.

## 2. Loss Calculation:

- The predicted output is compared to the actual target value for the input data point using a loss function (L). The loss function quantifies the error of the network's prediction. Common loss functions include Mean Squared Error (MSE) for regression and Cross-Entropy for classification.

### 3. Weight and Bias Update:

- Once the gradients of the loss with respect to all the weights and biases in the network are calculated, these gradients are used to update the weights and biases.
- The goal is to adjust the weights and biases in a direction that reduces the loss. This is typically done using an optimization algorithm, most commonly **Gradient Descent**.
- The update rule for a weight  $w_{ij}$  is:

$$w_{ij}^{new} = w_{ij}^{old} - \eta \frac{\partial L}{\partial w_{ij}}$$

- The update rule for a bias  $b_j$  is:

$$b_j^{new} = b_j^{old} - \eta \frac{\partial L}{\partial b_j}$$

- where  $\eta$  is the **learning rate**, a hyperparameter that controls the step size of the updates.

### 4. Backward Pass (Backpropagation<sup>1</sup> of Error):

- This is the crucial step where the network learns. The error calculated in the previous step is propagated backward through the network, layer by layer.
  - **Calculate the gradient of the loss with respect to the output layer's activations:**
- For each neuron in the output layer, the partial derivative of the loss function with respect to its activation ( $\partial a_j / \partial L$ ) is calculated.
  - **Calculate the gradient of the loss with respect to the weights and biases of the output layer:**
- Using the chain rule of calculus, the gradient of the loss with respect to the weights connecting the last hidden layer to the output layer ( $\partial w_{ij} / \partial L$ ) and the biases of the output layer ( $\partial b_j / \partial L$ ) are calculated. These gradients indicate how much each weight and bias contributed to the overall error.

➤ **Propagate the gradients backward to the hidden layers:**

- For each hidden layer, the gradient of the loss with respect to the activations of the neurons in that layer ( $\partial a_i / \partial L$ ) is calculated based on the gradients from the subsequent layer and the weights connecting those layers.
- Again, using the chain rule, the gradients of the loss with respect to the weights ( $\partial w_{ki} / \partial L$ ) and biases ( $\partial b_i / \partial L$ ) of the current hidden layer are computed.

➤ This backward propagation of gradients continues until the input layer is reached.

**5. Iteration:**

- Steps 1-4 are repeated for multiple **epochs** (passes through the entire training dataset) or until a certain stopping criterion is met (e.g., the loss on a validation set stops decreasing significantly).

**In essence, backpropagation efficiently computes how each weight and bias in the network affects the final error. By using these error signals (gradients), the optimization algorithm can iteratively adjust the network's parameters to learn the underlying patterns in the training data and improve its predictive performance.**

**Key Concepts in Backpropagation:**

- **Chain Rule:** The fundamental calculus rule that allows the calculation of gradients through composite functions (like the layers of a neural network).
- **Gradients:** Indicate the direction of the steepest increase in the loss function. By moving in the opposite direction of the gradient, we aim to minimize the loss.
- **Learning Rate:** A crucial hyperparameter that determines the size of the steps taken during weight and bias updates. A too-high learning rate can lead to instability, while a too-low learning rate can result in slow convergence.
- **Optimization Algorithms:** Algorithms like Stochastic Gradient Descent (SGD), Adam, RMSprop, etc., use the gradients calculated by backpropagation to efficiently update the network's parameters.